

LLVM DISTRIBUTORS CONFERENCE 2021

Thomas Preud'homme



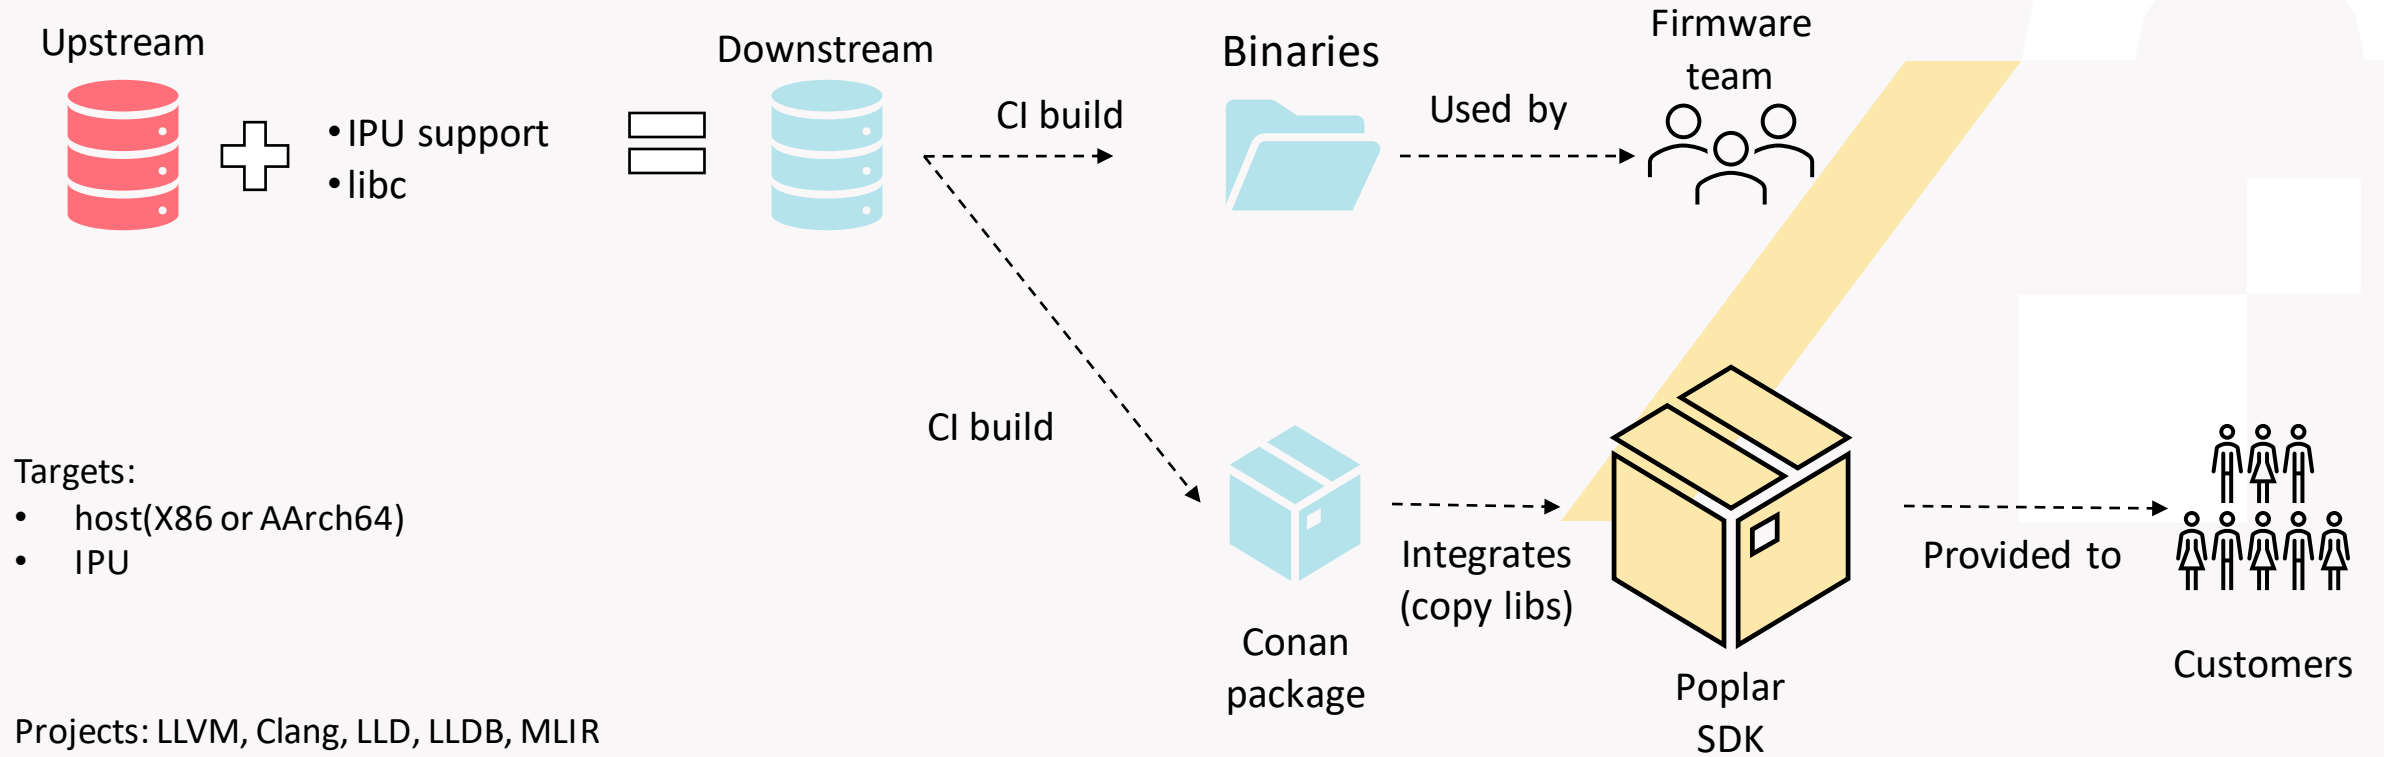
GRAPHCORE

Graphcore IPU: quick facts

- Designed to accelerate machine intelligence
- Contains 1472 cores running 6 parallel threads each
- Each core have their own fast local memory for a combined total of 900MB
- Development for IPU using Poplar® SDK
 - uses LLVM to compile for individual cores
 - uses MLIR for some of the high-level optimisation
- More info at <https://www.graphcore.ai/products/ipu>



Graphcore IPU LLVM distribution



External changes

Merges

- Target-specific code in shared files
- Merge process reviewability
- SDAG unit testing
- Behaviour changes in tools
- Ongoing freestanding proposal
- Running tests as root

Fast changing ISA

- Instruction tablegen generation

Target-specific code in shared files

```
TargetInfo *AllocateTarget(const llvm::Triple &Triple,  
                           const TargetOptions &Opts) {  
    llvm::Triple::OSType os = Triple.getOS();  
    switch (Triple.getArch()) {
```

Plenty of examples:

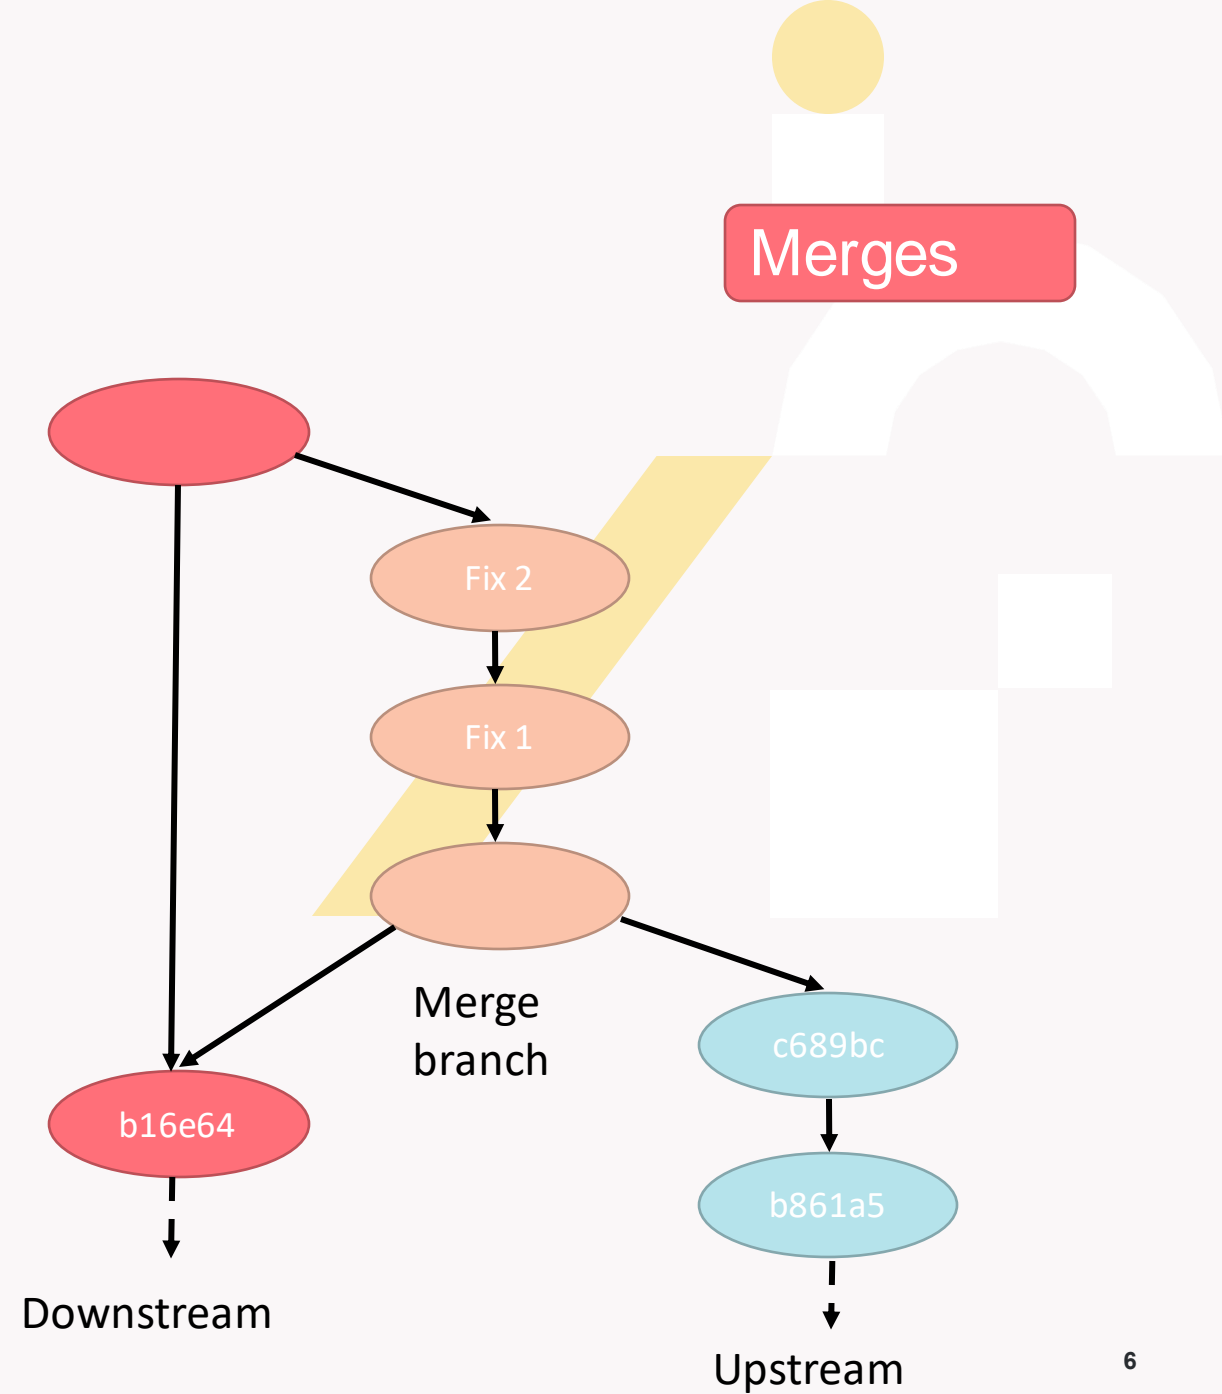
```
clang/include/clang/Basic/Attr.td  
clang/include/clang/Basic/AttrDocs.td  
clang/include/clang/Basic/TargetBuiltins.h  
clang/include/clang/Driver/Options.td  
clang/lib/Basic/Targets.cpp  
clang/lib/CodeGen/CGBuiltin.cpp  
clang/lib/CodeGen/TargetInfo.cpp  
clang/lib/Driver/CMakeLists.txt  
clang/lib/Driver/Driver.cpp  
clang/lib/Driver/ToolChains/Clang.cpp  
Clang/lib/Driver/ToolChains/CommonArgs.cpp  
Clang/lib/Sema/SemaDeclAttr.cpp  
(...)
```



Merges

Merge process reviewability

How to review merge-related changes?



SDAG unit testing

Motivation: SDAG node with no natural IR representation


Example: ANY_EXTEND only appear through combine

Solution: SDAG unit testing

```
; RUN: llvm-link %isdopc %s | opt -instcombine -always-inline | llc
```

```
@ISD_ANY_EXTEND = external constant i32  
declare i2 @llvm.ipu.SDAG.unary.i4.i2(i32, i2)
```

```
define i4 @test(i2 %x) {  
  %id = load i32, i32* @ISD_ANY_EXTEND  
  %res = call i4 @llvm.ipu.SDAG.unary.i4.i2(i32 %id, i2 %x)  
  ret i4 %res  
}
```



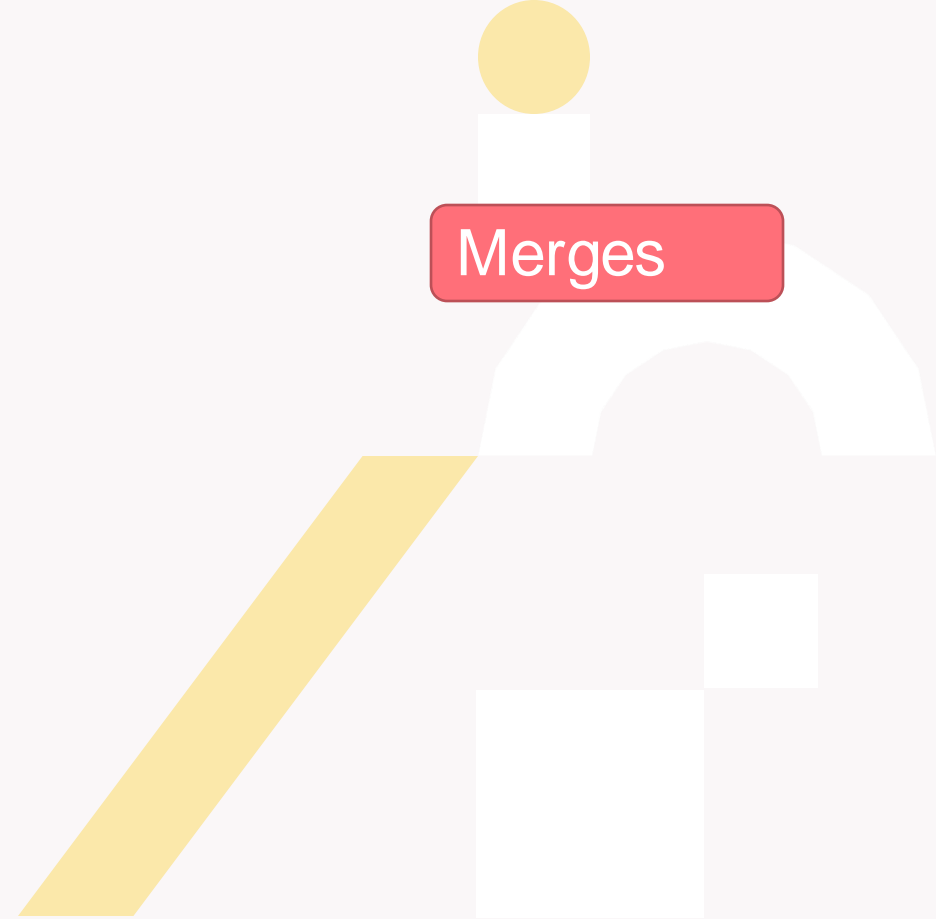
ISD constants file generation

ISDOpcodeConstantsPrinter.cpp:

```
const StringMap<std::pair<StringRef, int>> ISDNameConstantMapping = {  
    MAP(ISD_X, ISD::X),  
    (...)  
}  
  
for (auto &Mapping : ISDNameConstantMapping) {  
    std::tie(VarName, ISDValue) = Mapping.getValue();  
    std::cout << "extern constexpr unsigned " << VarName.str() << " = "  
        << ISDValue << ';' << std::endl;  
}
```

CMakeLists.txt:

1. Build ISD printer program with **host** compiler
2. Run program and compile output with target compiler



Behaviour changes in tools

Example: llvm-lit's --no-indirectly-run-check


Our setup: Lit invoked on individual tests by CMake in many repositories

Problem:

- Lit now errors out if a test would not have been run if invoked on parent directory
- Lit from both internal and external LLVM distribution

Solution 1: New commit to add lit config option to control it

Solution 2: Use lit from pip



Ongoing freestanding proposal


Context:

- no dynamic allocation on IPU for efficiency and robustness reasons
- current C++ specification for freestanding implementation requires dynamic allocation

Problem: Very invasive diff to add support

Our approach: copy libc++ in separate repo, allowing less frequent merge schedule

Note: current status of ongoing freestanding proposal no longer contains dynamic allocation




Running tests as root

Context: builds and tests run in docker as root

Problem: access right tests

Examples:

- `llvm/test/tools/llvm-ar/error-opening-permission.test`
- `llvm/test/tools/llvm-dwarfdump/X86/output.s`
- `llvm/test/tools/llvm-ifs/fail-file-write.test`



Tablegen generation

Motivation: Software-hardware co-design requires quick support of new instructions

Solution: Generation of instruction tablegen files & tests from ISA specification

```
class inst_x_aaan : Instruction, Sched<[Res]> {  
  
    let AsmString = "not $op0, $op1";  
    dag OutOperandList = (outs AR:$op0);  
    dag InOperandList = (ins AR:$op1);  
    (...)  
}
```

And if not manually defined:

```
def X : inst_x_aaan;
```

Fast changing ISA

Pros:

- Freely updated assembler
- Allow mix of manual and automated definitions

Cons:

- Not suitable for upstream



THANK YOU

Thomas Preud'homme
thomasp@graphcore.ai

